

Yuntian DENG

Google Scholar◊ Personal Website

dengyuntian@seas.harvard.edu◊ (412) 897-8303

Education

Harvard University	Cambridge, MA
Ph.D. in Computer Science	Expected May 2023
Advisors: Alexander Rush & Stuart Shieber	
Carnegie Mellon University	Pittsburgh, PA
Master in Language Technologies	Aug 2016
Advisor: Eric Xing	
Tsinghua University	Beijing, China
Bachelor of Engineering, Department of Automation	Jul 2014

Research Interests

My research aims at enabling ubiquitous long-form text generation, where anyone can produce a document on any topic with just a few clicks. To achieve this goal, my research focuses on the following three directions, utilizing advances in deep learning combined with probabilistic modeling:

Transparent Generation to enable users to understand, debug, and control document generation.

Long-Form Coherence to generate coherent documents with self-consistency and clear transitions.

Efficient Systems to handle the scale and real-time requirements of long-form text generation.

Awards

University of Chicago Rising Stars in Data Science	Oct 2022
ACM Gordon Bell Special Prize for Covid Research	Nov 2022
Nvidia Fellowship (10 awardees)	Dec 2021
Harvard Certificates of Distinction in Teaching	Spring 2019, Fall 2020, Fall 2021
ICML expert reviewer	Feb 2021
Twitch Fellowship Finalist	Dec 2020
NeurIPS 2020 top 10% reviewer	Oct 2020
DAC 2020 Best Paper Award	Jul 2020
Baidu Fellowship (10 awardees)	Feb 2019
French American Doctoral Exchange Program (Embassy of France, 10 US awardees)	Jul 2018

ACL 2017 Best Demo Paper Award Runner-up	Aug 2017
Evergrande Real Estate Group Scholarship	Nov 2013
Samsung SK Scholarship	Nov 2012
Chinese National Scholarship	Nov 2011

Publications

*Equal Contribution.

Selected Papers

- 1 **Yuntian Deng***, Yoon Kim*, Justin Chiu, Demi Guo, Alexander M. Rush. Latent Alignment and Variational Attention. *NeurIPS 2018*
- 2 **Yuntian Deng**, Anton Bakhtin, Myle Ott, Arthur Szlam, Marc'Aurelio Ranzato. Residual Energy-Based Models for Text Generation. *ICLR 2020*
- 3 **Yuntian Deng**, Alexander Rush. Cascaded Text Generation with Markov Transformers. *NeurIPS 2020*
- 4 **Yuntian Deng**, Volodymyr Kuleshov, Alexander Rush. Model Criticism for Long-Form Text Generation. *EMNLP 2022*
- 5 **Yuntian Deng**, Anssi Kanervisto, Jeffrey Ling, Alexander M. Rush. Image-to-Markup Generation with Coarse-to-Fine Attention. *ICML 2017* (Demo: im2markup.yuntiangdeng.com)
- 6 **Yuntian Deng**, Noriyuki Kojima, Alexander Rush. Markup-to-Image Diffusion Models with Scheduled Sampling. *ICLR 2023* (Demo: huggingface.co/spaces/yuntian-deng/latex2im)

Journal Papers

- 7 Anton Bakhtin*, **Yuntian Deng***, Sam Gross, Myle Ott, Marc'Aurelio Ranzato, Arthur Szlam. Residual Energy-Based Models for Text. *JMLR 2021*

Conference Papers

- 8 Richa Rastogi, Yair Schiff, Alon Hachohen, Zhaozhi Li, Ian Lee, **Yuntian Deng**, Mert R. Sabuncu, Volodymyr Kuleshov. Semi Parametric Inducing Point Networks. *ICLR 2023*
- 9 Maxim Zvyagin*, Alexander Brace*, Kyle Hippe*, **Yuntian Deng***, Bin Zhang, Cindy Orozco Bohorquez, Austin Clyde, Bharat Kale, Danilo Perez-Rivera, Heng Ma, Carla M. Mann, Michael Irvin, J. Gregory Pauloski, Logan Ward, Valerie Hayot, Murali Emani, Sam Foreman, Zhen Xie, Dian-gen Lin, Maulik Shukla, Weili Nie, Josh Romero, Christian Dallago, Arash Vahdat, Chaowei Xiao, Thomas Gibbs, Ian Foster, James J. Davis, Michael E. Papka, Thomas Brettin, Rick Stevens, Anima Anandkumar, Venkatram Vishwanath, Arvind Ramanathan, GenSLMs: Genome-scale language models reveal SARS-CoV-2 evolutionary dynamics. **ACM Gordon Bell Special Covid Prize**

- 10 Justin Chiu*, **Yuntian Deng***, Alexander Rush. Low-Rank Constraints for Fast Inference in Structured Models. *NeurIPS 2021*
- 11 **Yuntian Deng**, Alexander Rush. Sequence-to-Lattice Models for Fast Translation. *EMNLP 2021 Findings*
- 12 Keyon Vafa, **Yuntian Deng**, David Blei, Alexander Rush. Rationales for Sequential Predictions. *EMNLP 2021 oral*
- 13 Anton Bakhtin, Sam Gross, Myle Ott, **Yuntian Deng**, Marc’Aurelio Ranzato, Arthur Szlam. Real or Fake? Learning to Discriminate Machine from Human Generated Text. *arXiv 1906.03351*
- 14 Thierry Tambe, En-Yu Yang, Zishen Wan, **Yuntian Deng**, Vijay Janapa Reddi, Alexander Rush, David Brooks, Gu-Yeon Wei. Algorithm-Hardware Co-Design of Adaptive Floating-Point Encodings for Resilient Deep Learning Inference. *Design Automation Conference (DAC) 2020 **Best Paper Award***
- 15 Zachary Ziegler*, **Yuntian Deng***, Alexander Rush. Neural Linguistic Steganography. *EMNLP 2019 oral (Demo: steganography.live)*
- 16 **Yuntian Deng**, David Rosenberg, Gideon Mann. Challenges in end-to-end neural scientific table recognition. *ICDAR 2019*
- 17 Ran Zhao, **Yuntian Deng**, Mark Dredze, Arun Verma, David Rosenberg, Amanda Stent. Visual Attention Model for Cross-sectional Stock Return Prediction and End-to-End Multimodal Market Representation Learning. *FLAIRS 2019*
- 18 Sebastian Gehrmann, **Yuntian Deng**, Alexander M. Rush. Bottom-Up Abstractive Summarization. *EMNLP 2018*
- 19 Pengtao Xie, **Yuntian Deng**, Yi Zhou, Abhimanu Kumar, Yaoliang Yu, James Zou, Eric P Xing. Learning Latent Space Models with Angular Constraints. *ICML 2017*
- 20 Guillaume Klein, Yoon Kim, **Yuntian Deng**, Jean Senellart, Alexander M Rush. OpenNMT: Open-Source Toolkit for Neural Machine Translation. *ACL 2017 **Best Demo Paper Award Runner-Up***
- 21 Zichao Yang, Zhiting Hu, **Yuntian Deng**, Chris Dyer, Alex Smola. Neural Machine Translation with Recurrent Attention Modeling. *EACL 2017*
- 22 Xuezhe Ma, Yingkai Gao, Zhiting Hu, Yaoliang Yu, **Yuntian Deng**, Eduard Hovy. Dropout with expectation-linear regularization. *ICLR 2017*
- 23 Hao Zhang, Zhiting Hu, **Yuntian Deng**, Mrinmaya Sachan, Zhicheng Yan, Eric P. Xing. Learning Concept Taxonomies from Multi-modal Data. *ACL 2016*
- 24 Pengtao Xie, **Yuntian Deng**, Eric P. Xing. Diversifying Restricted Boltzmann Machine for Document Modeling. *KDD 2015*

25 Zhiting Hu, Poyao Huang, **Yuntian Deng**, Yingkai Gao, Eric P. Xing. Entity Hierarchy Embedding. *ACL 2015*

Teaching Experience

Harvard University

Head TA, Intro to Computational Linguistics and NLP (Stuart Shieber)	Aug 2021 - Dec 2021
Head TA, Intro to Computational Linguistics and NLP (Stuart Shieber)	Aug 2020 - Dec 2020
TA, Natural Language Processing (Alexander Rush)	Jan 2019 - May 2019
Lecture, Natural Language Processing - Translation (Alexander Rush)	Feb 13, 2019
TA, Advanced Machine Learning (Alexander Rush)	Aug 2017 - Dec 2017

Cornell University

Lab Material Preparation, Break Through Tech AI (Alexander Rush)	Aug 2021
--	----------

Carnegie Mellon University

TA, Graduate Probabilistic Graphical Models (Matthew Gormley, Eric Xing)	Jan 2016 - May 2016
TA, Graduate Machine Learning (Ziv Bar Joseph, Eric Xing)	Aug 2015 - Dec 2015

Revere High School

Volunteer Instructor of AP CS (Microsoft Philanthropies)	Sep 2020 - May 2021
--	---------------------

Professional Service

2022: Reviewer of ACL ARR, EMNLP, ICML, NeurIPS, ICLR, EACL, ML Reproducibility Challenge, ACL Demo, NAACL Demo, and EMNLP GEM Workshop.

2021: Expert Reviewer of ICML. Oral Session Volunteer of ACL. Reviewer of NeurIPS, ACL, ACL ARR, ICLR, AAAI, EACL, NAACL, ML Reproducibility Challenge, ICLR EBM Workshop, IEEE TIIS, and NEJLT. Volunteer of EMNLP.

2020: Top 10% Reviewer of NeurIPS. Reviewer of ICML, ACL, EMNLP, AACL, AAAI, ACL Demo, COLING, IEEE TIFS, Transactions on Information Systems, and Journal of Computer Science and Technology. Volunteer of ICML. Volunteer moderator of ACL.

2019: Reviewer of NeurIPS, EMNLP, ICLR, NAACL, AAAI, IEEE TNNLS, ACM Computing Surveys, NeurIPS LIRE Workshop, NeurIPS Reproducibility Challenge, EMNLP Summarization Workshop, ICML GraphReason Workshop, and ICLR DeepGenStruct Workshop.

2018: Reviewer of ICLR.

2017: Reviewer of ICML TADGM workshop.

2016: Volunteer of ICML.

Community Service

Organizer of Efficient NLP Workshop (in submission)	2022
Childcare Chair at NAACL 2022 D&I Committee	2022
Mentor of Harvard Women in CS (in charge of a bi-weekly reading group)	Sep 2020 - Present
Volunteer Instructor of AP CS at Revere High School	Sep 2020 - May 2021
Leader of Harvard English Language Table	Sep 2021 - Present

Internships

Nvidia / Argonne National Lab (Anima Anandkumar / Arvind Ramanathan)	May - Dec 2022
Facebook AI Research (Marc'Aurelio Ranzato)	May - Dec 2019
Bloomberg CTO Office (David Rosenberg, Gideon Mann)	Jan - Aug 2017
UCSD (Charles Elkan)	Jul - Sep 2013

Talks

U Chicago Rising Star: Model Criticism for Long-Form Text Generation	Nov 2022, Chicago
Princeton NLP: Model Criticism for Long-Form Text Generation	Nov 2022, Princeton
EMNLP 2021 Findings: Sequence-to-Lattice Models for Fast Translation	Oct 2021, remote
OpenAI: Residual Energy-based Models for Text Generation	Apr 2021, remote
NeurIPS 2020: Cascaded Text Generation with Markov Transformers	Oct 2020, remote
Baidu: Cascaded Text Generation with Markov Transformers	Jun 2020, remote
ICLR 2020: Residual Energy-based Models for Text Generation	Apr 2020, remote
LinkedIn: Residual Energy-based Models for Text Generation	Mar 2020, remote
Wayfair: Neural Encoder-Decoder Models	Nov 2019, Boston
EMNLP 2019: Neural Linguistic Steganography	Nov 2019, Hong Kong
FAIR NLP Meeting: Energy-Based Models for Text Generation	Aug 2019, New York City
NeurIPS 2018: Variational Attention	Dec 2018, Montreal, Canada

The French National Center for Scientific Research: Variational Attention	Jul 2018, France
Association for Machine Translation in the Americas: OpenNMT	Mar 2018, Boston
Open-Source Neural Machine Translation Workshop: Image-to-Text	Mar 2018, Paris, France
Tencent Social Network Group TSAIC: OpenNMT	Dec 2017, Shenzhen, China
ICML: Image-to-Markup Generation	Aug 2017, Sydney, Australia

Open Source Projects

Cascaded Generation (120 Github stars)	PyTorch
Neural Linguistic Steganography (157 Github stars)	PyTorch
Variational Attention (322 Github stars)	PyTorch
Image-to-Markup Generation (1k Github stars)	LuaTorch
OpenNMT (5.8k Github stars)	PyTorch
Attention OCR (1.1k Github stars)	Tensorflow

Mentees

Current

Richa Rastogi (Ph.D. student, Cornell Tech)
Celine Lee (Ph.D. student, Cornell Tech)

Past

Jambay Kinley (Undergraduate student, Harvard University → Microsoft)
Demi Guo (Undergraduate student, Harvard University → Ph.D. student, Stanford University)
Jeffrey Ling (Undergraduate student, Harvard University → Google → Startup)

References

Prof. Alexander Rush (Cornell University): arush@cornell.edu
Prof. Stuart Shieber (Harvard University): shieber@g.harvard.edu
Dr. Marc'Aurelio (DeepMind): ranzatomr@gmail.com
Prof. Volodymyr Kuleshov (Cornell University): vk379@cornell.edu